



# I am as a shadow product manager building this with complete help from you as AI expert, i am on day 10, describe each step for day 10 activities and help decode how to do it, in easy and elaborative step description

Based on the document, **Day 10** focuses on **Audit Logging & Encryption Hygiene**. Let me break down each step in easy, elaborative detail to help you implement this as a shadow product manager with AI assistance.<sup>[1]</sup>

## Day 10: Audit Logging & Encryption Hygiene

### Objective

Log all user actions in a secure, tamper-evident manner and verify that all sensitive data (PHI) is properly encrypted at rest.<sup>[1]</sup>

### Step 1: Create the Audit Logging Module ([audit.py](#))

**What you're building:** A system that records every action users take—like uploading notes, generating summaries, viewing logs—in a way that cannot be secretly modified later.<sup>[1]</sup>

#### How it works:

- Every time someone does something in your app, you'll write a log entry to a file
- Each log entry contains: timestamp, username, what action they performed, what data they touched, and a **hash** (like a digital fingerprint) of the previous log entry
- By linking each entry to the previous one through hashes, you create a **chain**—if someone tries to delete or edit an old entry, the chain breaks and you'll know tampering occurred<sup>[1]</sup>

**Create a new file:** `app/audit.py`

#### Code template:

```
import json
import hashlib
from datetime import datetime
from pathlib import Path
```

```

class AuditLogger:
    def __init__(self, log_file_path="logs/app_audit.jsonl"):
        self.log_file = Path(log_file_path)
        self.log_file.parent.mkdir(parents=True, exist_ok=True)
        # Create file if it doesn't exist
        if not self.log_file.exists():
            self.log_file.touch()

    def _get_last_hash(self):
        """Read the last log entry and return its hash"""
        try:
            with open(self.log_file, 'r') as f:
                lines = f.readlines()
                if lines:
                    last_entry = json.loads(lines[-1])
                    return last_entry.get('sha256_curr', '')
        except:
            pass
        return '' # First entry has no previous hash

    def _compute_hash(self, log_entry):
        """Create a hash fingerprint of the log entry"""
        # Convert the log entry to a string and hash it
        entry_string = json.dumps(log_entry, sort_keys=True)
        return hashlib.sha256(entry_string.encode()).hexdigest()

    def log_action(self, user, action, resource, additional_info=None):
        """
        Main logging function - call this whenever a user does something

        Args:
            user: username (e.g., 'dr_smith')
            action: what they did (e.g., 'UPLOAD_NOTE', 'GENERATE_SUMMARY', 'VIEW_LOGS')
            resource: what they acted on (e.g., 'note_12345.txt', 'patient_record')
            additional_info: any extra details (dictionary)
        """
        # Get the hash of the previous log entry
        previous_hash = self._get_last_hash()

        # Create the new log entry
        log_entry = {
            "timestamp": datetime.utcnow().isoformat() + "Z",
            "user": user,
            "action": action,
            "resource": resource,
            "sha256_prev": previous_hash,
            "additional_info": additional_info or {}
        }

        # Compute hash of THIS entry
        current_hash = self._compute_hash(log_entry)
        log_entry["sha256_curr"] = current_hash

        # Append to log file (append-only = cannot change old entries)
        with open(self.log_file, 'a') as f:
            f.write(json.dumps(log_entry) + '\n')

```

```
return log_entry
```

### What each part does:

- `__init__`: Sets up where logs will be stored (`logs/app_audit.jsonl`)
- `_get_last_hash`: Reads the last log entry to get its hash (needed for chaining)
- `_compute_hash`: Creates a unique fingerprint (SHA-256 hash) for each log entry
- `log_action`: The main function you'll call from your Streamlit app whenever something happens

## Step 2: Add OpenTelemetry-Style Attributes

**What this means:** OpenTelemetry is an industry-standard way of structuring logs so they can be exported to monitoring tools like OpenSearch later. You'll add special fields that make your logs compatible with this standard.<sup>[1]</sup>

**Update the `log_action` function** to include OpenTelemetry attributes:

```
import uuid

def log_action(self, user, action, resource, additional_info=None):
    """Enhanced with OpenTelemetry-style attributes"""
    previous_hash = self._get_last_hash()

    # Generate unique IDs for tracing
    trace_id = str(uuid.uuid4())
    span_id = str(uuid.uuid4())[:16] # Shorter ID for span

    log_entry = {
        "timestamp": datetime.utcnow().isoformat() + "Z",
        "user": user,
        "action": action,
        "resource": resource,
        "sha256_prev": previous_hash,
        "additional_info": additional_info or {},

        # OpenTelemetry attributes
        "otel_trace_id": trace_id,
        "otel_span_id": span_id,
        "otel_service_name": "clinical-rag-app",
        "severity": "INFO" # Can be DEBUG, INFO, WARN, ERROR
    }

    current_hash = self._compute_hash(log_entry)
    log_entry["sha256_curr"] = current_hash

    with open(self.log_file, 'a') as f:
        f.write(json.dumps(log_entry) + '\n')
```

```
return log_entry
```

### Why these fields matter:

- `trace_id`: Unique identifier that can link related actions together
- `span_id`: Sub-identifier for individual operations within a trace
- `service_name`: Identifies which application created the log
- `severity`: Helps filter logs by importance level

## Step 3: Integrate Audit Logging into Your Streamlit App

**Where to add logging:** Anywhere users interact with sensitive data.<sup>[1]</sup>

**Update your `app/main.py` (Streamlit app):**

```
import streamlit as st
from audit import AuditLogger

# Initialize the audit logger at the top
audit_logger = AuditLogger()

# Example: Log when user uploads a note
if st.button("Upload Note"):
    # Your existing upload code...

    # Add audit log
    audit_logger.log_action(
        user=st.session_state.get('username', 'anonymous'),
        action="UPLOAD_NOTE",
        resource=uploaded_file.name,
        additional_info={"file_size": len(note_text)}
    )

# Example: Log when user generates a summary
if st.button("Generate Summary"):
    # Your existing summarization code...

    audit_logger.log_action(
        user=st.session_state.get('username', 'anonymous'),
        action="GENERATE_SUMMARY",
        resource=f"note_{note_id}",
        additional_info={"retrieval_count": len(retrieved_chunks)}
    )

# Example: Log when admin views logs
if user_role == "admin" and st.button("View Audit Logs"):
    audit_logger.log_action(
        user=st.session_state['username'],
        action="VIEW_LOGS",
        resource="app_audit.jsonl"
    )
```

```
# Display logs
with open("logs/app_audit.jsonl", "r") as f:
    st.text(f.read())
```

### Key actions to log:

- UPLOAD\_NOTE: When notes are uploaded
- GENERATE\_SUMMARY: When AI generates summaries
- VIEW\_LOGS: When someone accesses audit logs
- ACCESS\_PHI: If anyone accesses encrypted PHI data
- DEID\_PROCESS: When de-identification runs
- INDEX\_UPDATE: When vector database is updated

## Step 4: Verify Encryption of PHI (Fernet Encryption)

**What you're checking:** All PHI (Protected Health Information) must be encrypted before being saved to disk.<sup>[1]</sup>

**Review your existing** `deid_pipeline.py` (from Day 5) to confirm encryption is working:

```
from cryptography.fernet import Fernet
import json
from pathlib import Path

# Load or generate encryption key (DO NOT store in Git)
KEY_FILE = Path("secure_store/encryption.key")
KEY_FILE.parent.mkdir(parents=True, exist_ok=True)

if not KEY_FILE.exists():
    key = Fernet.generate_key()
    with open(KEY_FILE, 'wb') as f:
        f.write(key)
else:
    with open(KEY_FILE, 'rb') as f:
        key = f.read()

cipher = Fernet(key)

# Example: Encrypting PHI span map
def save_phi_map(note_id, phi_spans):
    """
    Encrypt and save the mapping of PHI locations

    Args:
        note_id: unique identifier for the note
        phi_spans: dictionary mapping PHI entities to their original values
    """
    phi_json = json.dumps(phi_spans)
    encrypted_data = cipher.encrypt(phi_json.encode())
```

```

# Save encrypted data
encrypted_file = Path(f"secure_store/{note_id}_phi.enc")
with open(encrypted_file, 'wb') as f:
    f.write(encrypted_data)

print(f"✓ PHI map encrypted and saved to {encrypted_file}")

# Example: Decrypting PHI (only for authorized clinicians)
def load_phi_map(note_id):
    """Decrypt and load PHI map"""
    encrypted_file = Path(f"secure_store/{note_id}_phi.enc")

    with open(encrypted_file, 'rb') as f:
        encrypted_data = f.read()

    decrypted_json = cipher.decrypt(encrypted_data).decode()
    return json.loads(decrypted_json)

```

### Verification checklist:

- ✓ Encryption key stored in `secure_store/encryption.key`
- ✓ Key file added to `.gitignore` (never committed to GitHub)
- ✓ All PHI files have `.enc` extension
- ✓ Original PHI text is never saved unencrypted

## Step 5: Implement Key Rotation with MultiFernet

**What is key rotation?** Periodically changing encryption keys to improve security. If one key is compromised, older data remains protected by previous keys.<sup>[1]</sup>

**Create** `app/key_rotation.py`:

```

from cryptography.fernet import Fernet, MultiFernet
import json
from pathlib import Path
from datetime import datetime

class KeyManager:
    def __init__(self, key_dir="secure_store/keys"):
        self.key_dir = Path(key_dir)
        self.key_dir.mkdir(parents=True, exist_ok=True)
        self.keys_file = self.key_dir / "key_history.json"

        # Initialize or load keys
        if not self.keys_file.exists():
            self._create_initial_key()

        self.load_keys()

    def _create_initial_key(self):

```

```

    """Generate the first encryption key"""
    key = Fernet.generate_key()
    key_record = {
        "keys": [
            {
                "key": key.decode(),
                "created_at": datetime.utcnow().isoformat() + "Z",
                "active": True
            }
        ]
    }
    with open(self.keys_file, 'w') as f:
        json.dump(key_record, f, indent=2)

def load_keys(self):
    """Load all keys (current + historical)"""
    with open(self.keys_file, 'r') as f:
        data = json.load(f)

    # Convert string keys to Fernet objects
    fernet_keys = [Fernet(k['key'].encode()) for k in data['keys']]

    # MultiFernet tries each key in order for decryption
    self.cipher = MultiFernet(fernet_keys)
    self.current_key = fernet_keys[0] # Most recent key

def rotate_key(self):
    """Generate a new key and add it to the front"""
    new_key = Fernet.generate_key()

    with open(self.keys_file, 'r') as f:
        data = json.load(f)

    # Add new key at the beginning
    data['keys'].insert(0, {
        "key": new_key.decode(),
        "created_at": datetime.utcnow().isoformat() + "Z",
        "active": True
    })

    # Mark old keys as inactive
    for k in data['keys'][1:]:
        k['active'] = False

    with open(self.keys_file, 'w') as f:
        json.dump(data, f, indent=2)

    print(f"✓ New key generated. Total keys in rotation: {len(data['keys'])}")
    self.load_keys() # Reload

def encrypt(self, plaintext):
    """Encrypt using the current (newest) key"""
    return self.cipher.encrypt(plaintext.encode())

def decrypt(self, ciphertext):
    """Decrypt using any available key (tries newest first)"""

```

```

        return self.cipher.decrypt(ciphertext).decode()

# Usage example
if __name__ == "__main__":
    km = KeyManager()

    # Encrypt some PHI
    encrypted = km.encrypt("Patient: John Doe, MRN: 12345")
    print(f"Encrypted: {encrypted[:50]}...")

    # Rotate key
    km.rotate_key()

    # Old encrypted data still decrypts with old key
    decrypted = km.decrypt(encrypted)
    print(f"Decrypted after rotation: {decrypted}")

```

### How MultiFernet works:

- Stores multiple keys in a list (newest first)
- Encrypts with the **newest** key
- Decrypts by trying each key until one works (automatic backward compatibility)

## Step 6: Document Key Rotation Policy

Create docs/key\_rotation\_policy.md:

```

# Encryption Key Rotation Policy

## Overview
All PHI is encrypted at rest using Fernet symmetric encryption. Keys are rotated quarterly.

## Key Management
- **Storage:** Keys stored in `secure_store/keys/` (excluded from Git)
- **Algorithm:** Fernet (AES-128 in CBC mode with HMAC authentication)
- **Rotation:** MultiFernet allows multiple keys for backward compatibility

## Rotation Schedule
- **Frequency:** Every 90 days
- **Process:**
    1. Generate new key using `KeyManager.rotate_key()`
    2. New key becomes active for all new encryption
    3. Historical keys retained for decryption of old data
    4. After 2 years, consider re-encrypting old data with current key

## Manual Rotation Command

```

python app/key\_rotation.py

```

## Security Controls

```



- Keys never committed to version control (`.gitignore`)
- Key file permissions set to 600 (owner read/write only)
- Audit log entry created on each rotation

## Step 7: Test Everything Together

Create a test script `app/test_day10.py`:

```
from audit import AuditLogger
from key_rotation import KeyManager
import json

def test_audit_logging():
    """Test that audit logs are created correctly"""
    print("=== Testing Audit Logger ===")
    logger = AuditLogger()

    # Simulate user actions
    logger.log_action("dr_smith", "UPLOAD_NOTE", "note_001.txt", {"size": 1024})
    logger.log_action("dr_smith", "GENERATE_SUMMARY", "note_001.txt")
    logger.log_action("admin", "VIEW_LOGS", "app_audit.jsonl")

    # Verify hash chain
    with open("logs/app_audit.jsonl", "r") as f:
        lines = f.readlines()
        for i, line in enumerate(lines[-3:]):
            entry = json.loads(line)
            print(f"Entry {i+1}: {entry['action']} by {entry['user']}")
            print(f"  Previous hash: {entry['sha256_prev'][:16]}...")
            print(f"  Current hash: {entry['sha256_curr'][:16]}...")

    print("✓ Audit logging test complete\n")

def test_encryption():
    """Test encryption and key rotation"""
    print("=== Testing Encryption & Key Rotation ===")
    km = KeyManager()

    # Encrypt PHI
    original = "Patient: Jane Doe, DOB: 1980-05-15, SSN: 123-45-6789"
    encrypted = km.encrypt(original)
    print(f"Original: {original}")
    print(f"Encrypted: {encrypted[:50]}...")

    # Decrypt
    decrypted = km.decrypt(encrypted)
    assert decrypted == original, "Decryption failed!"
    print(f"Decrypted: {decrypted}")

    # Rotate key
    print("\n--- Rotating key ---")
    km.rotate_key()

    # Old data still decrypts
```

```

    decrypted_after_rotation = km.decrypt(encrypted)
    assert decrypted_after_rotation == original, "Old data can't be decrypted!"
    print(f"✓ Old encrypted data still decrypts: {decrypted_after_rotation}")

    # New encryption uses new key
    new_encrypted = km.encrypt("New patient data")
    print(f"✓ New encryption successful: {new_encrypted[:50]}...")

    print("\n Encryption & rotation test complete\n")

if __name__ == "__main__":
    test_audit_logging()
    test_encryption()
    print("=== All Day 10 Tests Passed! ===")

```

### Run the test:

```

cd app
python test_day10.py

```

### Expected output:

```

=== Testing Audit Logger ===
Entry 1: UPLOAD_NOTE by dr_smith
    Previous hash: ...
    Current hash: ...
Entry 2: GENERATE_SUMMARY by dr_smith
    Previous hash: [matches previous current hash]
    Current hash: ...
Entry 3: VIEW_LOGS by admin
✓ Audit logging test complete

=== Testing Encryption & Key Rotation ===
Original: Patient: Jane Doe...
Encrypted: gAAAAABm...
Decrypted: Patient: Jane Doe...

--- Rotating key ---
✓ New key generated. Total keys in rotation: 2
✓ Old encrypted data still decrypts: Patient: Jane Doe...
✓ New encryption successful: gAAAAABm...
✓ Encryption & rotation test complete

=== All Day 10 Tests Passed! ===

```

## Step 8: Update Your .gitignore

Add these lines to .gitignore to ensure keys never get committed:

```
# Encryption keys - NEVER commit these!
secure_store/
*.key
*_phi.enc

# Audit logs (optional - some teams commit logs, others don't)
logs/
*.jsonl
```

## Deliverables for Day 10

By the end of today, you should have:

1. ✓ app/audit.py: Audit logger with hash-chained, append-only logs<sup>[1]</sup>
2. ✓ **OpenTelemetry attributes**: trace\_id, span\_id, service\_name in logs<sup>[1]</sup>
3. ✓ logs/app\_audit.jsonl: Sample audit log file with real entries<sup>[1]</sup>
4. ✓ app/key\_rotation.py: Key manager with MultiFernet rotation<sup>[1]</sup>
5. ✓ docs/key\_rotation\_policy.md: Documentation of rotation schedule and procedures<sup>[1]</sup>
6. ✓ **Updated** app/main.py: Streamlit app with audit logging on all critical actions<sup>[1]</sup>
7. ✓ .gitignore: Ensures keys and encrypted PHI never get pushed to GitHub<sup>[1]</sup>
8. ✓ **Test script**: Validates audit logging and encryption work correctly<sup>[1]</sup>

## Why This Matters for Interviews

When discussing this portfolio project with hiring managers, Day 10 demonstrates:

- **Security-first mindset**: You understand that logging and encryption aren't optional—they're core requirements for healthcare applications<sup>[1]</sup>
- **Compliance awareness**: Audit logs and encryption directly map to HIPAA Technical Safeguards (audit controls, integrity, person/entity authentication)<sup>[1]</sup>
- **Tamper evidence**: Hash-chained logs show sophisticated understanding beyond basic logging<sup>[1]</sup>
- **Operational readiness**: Key rotation policy shows you think about long-term security maintenance<sup>[1]</sup>
- **OpenTelemetry knowledge**: Using industry-standard log formats shows you're thinking about production observability<sup>[1]</sup>

This is exactly the kind of detail that separates a "built a demo" portfolio from a "thought through real-world deployment" portfolio.<sup>[1]</sup>



1. COMPLETE-WEEK-BY-WEEK-TASKS.docx